

Database Programming with SQL

9-1: Using GROUP BY and HAVING Clauses

Practice Activities

Objectives

- Construct and execute a SQL query using GROUP BY
- Construct and execute a SQL query using GROUP BY ... HAVING
- Construct and execute a GROUP BY on more than one column
- Nest group functions

Vocabulary

Identify the vocabulary word for each definition below.

HAVING clause	Used to specify which groups are to be displayed; restricts groups that do not meet group criteria
GROUP BY clause	Divides the rows in a table into groups

Try It / Solve It

1. In the SQL query shown below, which of the following is true about this query?

- F a. Kimberly Grant would not appear in the results set.
- F b. The GROUP BY clause has an error because the manager_id is not listed in the SELECT clause.
- T c. Only salaries greater than 16001 will be in the result set.
- T d. Names beginning with Ki will appear after names beginning with Ko.
- F e. Last names such as King and Kochhar will be returned even if they don't have salaries > 16000.

```
SELECT last_name, MAX(salary)
FROM employees
WHERE last_name LIKE 'K%'
GROUP BY manager_id, last_name
HAVING MAX(salary) >16000
ORDER BY last_name DESC ;
```

2. Each of the following SQL queries has an error. Find the error and correct it. Use Oracle Application Express to verify that your corrections produce the desired results.

- | | |
|--|---|
| a. SELECT manager_id
FROM employees
WHERE AVG(salary) <16000
GROUP BY manager_id; | SELECT manager_id
FROM employees
GROUP BY manager_id
HAVING AVG(salary) < 16000; |
| b. SELECT cd_number, COUNT(title)
FROM d_cds
WHERE cd_number < 93; | SELECT cd_number, COUNT(title)
FROM d_cds
WHERE cd_number < 93
GROUP BY cd_number; |
| c. SELECT ID, MAX(ID), artist AS Artist
FROM d_songs
WHERE duration IN('3 min', '6 min', '10 min')
HAVING ID < 50
GROUP by ID; | SELECT MAX(ID) AS ID, artist AS Artist
FROM d_songs
WHERE duration IN ('3 min', '6 min', '10 min')
GROUP BY artist
HAVING MAX(ID) < 50; |
| d. SELECT loc_type, rental_fee AS Fee
FROM d_venues
WHERE id <100
GROUP BY "Fee"
ORDER BY 2; | SELECT loc_type, rental_fee AS Fee
FROM d_venues
WHERE id < 100
GROUP BY loc_type, rental_fee
ORDER BY 2; |

3. Rewrite the following query to accomplish the same result:

SELECT DISTINCT MAX(song_id) FROM d_track_listings WHERE track IN (1, 2, 3);	SELECT MAX(song_id) FROM d_track_listings WHERE track IN (1, 2, 3);
---	---

4. Indicate True or False

- T a. If you include a group function and any other individual columns in a SELECT clause, then each individual column must also appear in the GROUP BY clause.
- F b. You can use a column alias in the GROUP BY clause.
- F c. The GROUP BY clause always includes a group function.

5. Write a query that will return both the maximum and minimum average salary grouped by department from the employees table.

```
SELECT department_id, MAX(avg_salary) AS max_avg_salary, MIN(avg_salary) AS min_avg_salary
FROM (
    SELECT department_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY department_id
)
GROUP BY department_id
ORDER BY 1;
```

6. Write a query that will return the average of the maximum salaries in each department for the employees table.

```
SELECT department_id, AVG(max_salary) AS avg_max_salary
FROM (
    SELECT department_id, MAX(salary) AS max_salary
    FROM employees
    GROUP BY department_id
)
GROUP BY department_id
ORDER BY 1;
```

Database Programming with SQL

9-2: Using ROLLUP and CUBE Operations and GROUPING SETS

Practice Activities

Objectives

- Use ROLLUP to produce subtotal values
- Use CUBE to produce cross-tabulation values
- Use GROUPING SETS to produce a single result set
- Use the GROUPING function to identify the extra row values created by either a ROLLUP or CUBE operation

Vocabulary

Identify the vocabulary word for each definition below.

ROLLUP	Used to create subtotals that roll up from the most detailed level to a grand total, following a grouping list specified in the clause
CUBE	An extension to the GROUP BY clause like ROLLUP that produces cross-tabulation reports
GROUPING SETS	Used to specify multiple groupings of data

Try It / Solve It

1. Within the Employees table, each manager_id is the manager of one or more employees who each have a job_id and earn a salary. For each manager, what is the total salary earned by all of the employees within each job_id? Write a query to display the Manager_id, job_id, and total salary. Include in the result the subtotal salary for each manager and a grand total of all salaries.

```
SELECT manager_id, job_id, SUM(salary) AS total_salary
FROM employees
GROUP BY ROLLUP (manager_id, job_id);
```

2. Amend the previous query to also include a subtotal salary for each job_id regardless of the manager_id.

```
SELECT manager_id, job_id, SUM(salary) AS total_salary
FROM employees
GROUP BY CUBE (manager_id, job_id);
```

3. Using GROUPING SETS, write a query to show the following groupings:

- department_id, manager_id, job_id
- manager_id, job_id
- department_id, manager_id

```
SELECT department_id, manager_id, job_id, SUM(salary) AS
total_salary
FROM employees
GROUP BY GROUPING SETS (
    (department_id, manager_id, job_id),
    (manager_id, job_id),
    (department_id, manager_id)
);
```

Database Programming with SQL

9-3: Set Operators

Practice Activities

Objectives

- Define and explain the purpose of SET operators
- Use a set operator to combine multiple queries into a single query
- Control the order of rows returned using set operators

Vocabulary

Identify the vocabulary word for each definition below.

UNION	operator that returns all rows from both tables and eliminates duplicates
NULL Columns	columns that were made up to match queries in another table that are not in both tables
UNION ALL	operator that returns all rows from both tables, including duplicates
SET Operators	used to combine results into one single result from multiple SELECT statements
MINUS	operator that returns rows that are unique to each table
INTERSECT	operator that returns rows common to both tables

Try It / Solve It

1. Name the different Set operators? **UNION, UNION ALL, INTERSECT, MINUS**
2. Write one query to return the employee_id, job_id, hire_date, and department_id of all employees and a second query listing employee_id, job_id, start_date, and department_id from the job_history table and combine the results as one single output. Make sure you suppress duplicates in the output.

```
SELECT employee_id, job_id, hire_date AS "DATE", department_id
FROM employees
UNION
SELECT employee_id, job_id, start_date AS "DATE", department_id
FROM job_history;
```
3. Amend the previous statement to not suppress duplicates and examine the output. How many extra rows did you get returned and which were they? Sort the output by employee_id to make it easier to spot.

```
SELECT employee_id, job_id, hire_date AS "DATE", department_id
FROM employees
UNION ALL
SELECT employee_id, job_id, start_date AS "DATE", department_id
FROM job_history
ORDER BY employee_id;
```
4. List all employees who have not changed jobs even once. (Such employees are not found in the job_history table)

```
SELECT employee_id, job_id, hire_date, department_id
FROM employees
WHERE employee_id NOT IN (SELECT employee_id FROM job_history);
```
5. List the employees that HAVE changed their jobs at least once.

```
SELECT employee_id, job_id, hire_date, department_id
FROM employees
WHERE employee_id IN (SELECT employee_id FROM job_history);
```
6. Using the UNION operator, write a query that displays the employee_id, job_id, and salary of ALL present and past employees. If a salary is not found, then just display a 0 (zero) in its place.

```
SELECT employee_id, job_id, salary
FROM employees
UNION
SELECT employee_id, job_id, 0 AS salary
FROM job_history;
```